

sentiment_analysis

July 26, 2019

1 Crisis Sentiment Analysis

This activity is a mini-project where students will create a data visualization dashboard, they have to analyze sentiment and tone about the news related to the financial crisis of 2008 that were published along the last month. Students will retrieve the news articles from the News API; by default, the developer account gives access to news articles up to a month old.

In this activity, students will use their new sentiment analysis skills, in combination to some of the skills they already master such as: Pandas, Pyviz, Plotly Express and PyViz Panel.

This Jupyter notebook is a sandbox where students will conduct the sentiment analysis tasks and charts creation before assembling the dashboard.

```
[1]: # Initial imports
import os
from path import Path
import pandas as pd
import numpy as np
import hvplot.pandas
import nltk
from wordcloud import WordCloud
from nltk.sentiment.vader import SentimentIntensityAnalyzer
from newsapi import NewsApiClient
from ibm_watson import ToneAnalyzerV3
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
import plotly.express as px
import matplotlib.pyplot as plt
import matplotlib as mpl
import panel as pn

plt.style.use("seaborn-whitegrid")
pn.extension("plotly")
```

1.1 Instructions

1.1.1 Fetching the Latests News Metions About the Crisis of 2008

Using the News API, get all the news in English about the financial crisis of 2008 using the keywords "financial AND crisis AND 2008" in the q parameter. Define a page_size=100 to have

at least 100 news articles to analyze.

```
[2]: # Retrieve the News API key
news_api = os.getenv("news_api")

[3]: # Create the newsapi client
newsapi = NewsApiClient(api_key=news_api)

[4]: # Fetch the news articles about the financial crisis on 2008 in English
crisis_news_en = newsapi.get_everything(
    q="financial AND crisis AND 2008", language="en", page_size=100
)

# Show the total number of news
crisis_news_en["totalResults"]
```

[4]: 2153

1.1.2 Creating a VADER Sentiment Scoring Function

Use the VADER sentiment scoring function from NLTK to score the sentiment polarity of the 100 news you fetched. Just for convenience, start downloading the `vader_lexicon` in order to initialize the VADER sentiment analyzer.

```
[5]: # Download/Update the VADER Lexicon
nltk.download("vader_lexicon")
```

```
[nltk_data] Downloading package vader_lexicon to
[nltk_data] /Users/josearturomorasoto/nltk_data...
[nltk_data] Package vader_lexicon is already up-to-date!
```

[5]: True

```
[6]: # Initialize the VADER sentiment analyzer
analyzer = SentimentIntensityAnalyzer()
```

In order to score the VADER sentiment, create a function named `get_sentiment_scores(text, date, source, url)` that will receive four parameters.

- `text` is the text whose sentiment will be scored.
- `date` the date the news article was published using the format `YYYY-MM-DD`.
- `source` is the name of the news article's source.
- `url` is the URL that points to the article.

The `get_sentiment_score()` function should return a Python dictionary with the scoring results. This dictionary is going to be used in the next section to create a `DataFrame`; the structure of the dictionary is the following:

- `date` the date passed as parameter to the function.
- `text` the text passed a parameter to the function.
- `source` the source passed as parameter to the function.

- url the URL passed as parameter to the function.
- compound the compound score from the VADER sentiment analyzer.
- pos the positive score from the VADER sentiment analyzer.
- neu the neutral score from the VADER sentiment analyzer.
- neg the negative score from the VADER sentiment analyzer.
- normalized the normalized scored based on the compound results. Its value should be 1 for positive sentiment, -1 for negative sentiment, and 0 for neutral sentiment.

This is an example of the function's return value:

```
{'date': '2019-06-24',
'text': '\nMore than a decade since the global economic meltdown of 2008
devastated lives across the world, no one who caused the crisis has
been held responsible.\n\n"The 2008 financial crisis displayed what
the world now identifies as financial contagion," says Philip J Baker,
the former managing partner of a US-based \nhedge fund that collapsed
during the financial crisis.\n\nDespite this, "zero Wall Street chief
executives have been to prison, even though there is today absolutely
no doubt that Wall Street executives and politicians \nwere complicit
in creating the crisis," he says. \n\nBaker was among the few
relatively smaller players imprisoned for the part they played.\n\n
In July 2009, he was arrested in Germany and extradited to the
United States where he faced federal court on charges of fraud and
financial crimes.\n\nHe pled guilty and was sentenced to 20 years
in prison for costing some 900 investors about $294mn worldwide.
He served eight years in jail and is now on \nparole and advocates
against financial crime.\n',
'source': 'aljazeera',
'url': 'https://www.aljazeera.com/programmes/specialseries/2019/06/men-stole-world-2008-financ',
'compound': -0.9911,
'pos': 0.048,
'neu': 0.699,
'neg': 0.254,
'normalized': -1}
```

```
[7]: # Define a function to get the sentiment scores
def get_sentiment_scores(text, date, source, url):
    sentiment_scores = {}

    # Sentiment scoring with VADER
    text_sentiment = analyzer.polarity_scores(text)
    sentiment_scores["date"] = date
    sentiment_scores["text"] = text
    sentiment_scores["source"] = source
    sentiment_scores["url"] = url
    sentiment_scores["compound"] = text_sentiment["compound"]
    sentiment_scores["pos"] = text_sentiment["pos"]
    sentiment_scores["neu"] = text_sentiment["neu"]
```

```

sentiment_scores["neg"] = text_sentiment["neg"]
if text_sentiment["compound"] >= 0.05: # Positive
    sentiment_scores["normalized"] = 1
elif text_sentiment["compound"] <= -0.05: # Negative
    sentiment_scores["normalized"] = -1
else:
    sentiment_scores["normalized"] = 0 # Neutral

return sentiment_scores

```

1.1.3 Creating the News Articles' Sentiments DataFrame

In this section you have to create a DataFrame that is going to be used to plot the sentiment analysis results. Using a for-loop, iterate across all the news articles you fetched to create the DataFrame structure; define an empty list to append the sentiment scoring results for each news article and create the DataFrame using the list as data source.

Once you create the DataFrame do the following:

- Sort the DataFrame rows by the date column.
- Define the date column as the DataFrame index.
- Save the DataFrame as a CSV file in order to use it on the sentiment analysis dashboard creation.

```

[8]: # Empty list to store the DataFrame structure
sentiments_data = []

# Loop through all the news articles
for article in crisis_news_en["articles"]:
    try:
        # Get sentiment scoring using the get_sentiment_score() function
        sentiments_data.append(
            get_sentiment_scores(
                article["content"],
                article["publishedAt"][:10],
                article["source"]["name"],
                article["url"],
            )
        )
    except AttributeError:
        pass

# Create a DataFrame with the news articles' data and their sentiment scoring_
→ results
crisis_news_df = pd.DataFrame(sentiments_data)

# Sort the DataFrame rows by date
crisis_news_df = crisis_news_df.sort_values(by="date")

```

```
# Define the date column as the DataFrame's index
crisis_news_df.set_index("date", inplace=True)
crisis_news_df.head()
```

```
[8]:
```

	compound	neg	neu	normalized	pos	source	\
date							
2019-06-25	-0.1027	0.038	0.962	-1	0.000	Forbes.com	
2019-06-25	-0.8625	0.206	0.794	-1	0.000	Politico	
2019-06-25	0.1027	0.055	0.881	1	0.064	Slate.com	
2019-06-26	-0.6705	0.111	0.889	-1	0.000	Fastcompany.com	
2019-06-26	0.0000	0.000	1.000	0	0.000	Digg.com	

	text	\
date		
2019-06-25	Share to facebookShare to twi...	
2019-06-25	Michael Kruse is a senior staff writer for Pol...	
2019-06-25	Andrew Yang didnt set out to become a test pre...	
2019-06-26	For the latter half of 2016, the Standing Rock...	
2019-06-26	We're halfway into 2019 and so far this year w...	

	url
date	
2019-06-25	https://www.forbes.com/sites/mayrarodriguezval...
2019-06-25	https://www.politico.com/magazine/story/2019/0...
2019-06-25	https://slate.com/news-and-politics/2019/06/an...
2019-06-26	https://www.fastcompany.com/90364616/public-ba...
2019-06-26	http://digg.com/2019/20-biggest-bankruptcies-c...

```
[9]: # Save the news articles DataFrame with VADER Sentiment scoring as a CSV file
file_path = Path("Data/news_vader.csv")
crisis_news_df.to_csv(file_path)
```

1.1.4 Creating the Average Sentiment Chart

Use hvPlot to create a two lines chart that compares the average compound and normalized sentiment scores along the last month.

```
[10]: avg_sent_data = (
    crisis_news_df[["compound", "normalized"]].groupby(by=crisis_news_df.index).
    ↪mean()
)
avg_sent_data.head()
```

```
[10]:
```

	compound	normalized
date		
2019-06-25	-0.287500	-0.333333
2019-06-26	-0.223500	-0.333333
2019-06-27	0.424033	0.666667

```
2019-06-28 -0.131675    -0.250000
2019-06-29  0.483250     1.000000
```

```
[11]: avg_sent_plot = avg_sent_data.hvplot(
        title="Average Sentiment About the Economic Crisis of 2008 Last Month",
        rot=90
    )

avg_sent_plot
```

```
[11]: :NdOverlay    [Variable]
      :Curve      [date]    (value)
```

1.1.5 Creating the Sentiment Distribution Chart

Based on the normalized sentiment score, create a bar chart using hvPlot that shows the number of negative, neutral and positive news articles. This chart represents the overall sentiment distribution.

```
[12]: sentiment_chart_df = (
        crisis_news_df[["normalized", "text"]].groupby("normalized").count()
    )
sentiment_chart_df.rename(
    index={-1: "Negative", 0: "Neutral", 1: "Positive"}, inplace=True
)
sentiment_chart_df
```

```
[12]:          text
normalized
Negative      43
Neutral       11
Positive      46
```

```
[13]: sentiment_bar_chart = sentiment_chart_df.hvplot.bar(
        xlabel="Sentiment",
        ylabel="Number of News",
        title="Overall Sentiment Distribution",
        color="text",
    )

sentiment_bar_chart
```

```
[13]: :Bars      [normalized]    (text)
```

1.1.6 Getting the Top 10 Positive and Negative News Articles

In this section you have to create two DataFrames, one with the top 10 positive news according to the compound score, and other with the top 10 negative news. Refer to the [hvplot.table\(\) documentation](#) to create two tables presenting the following columns of these news articles:

- Date

- Source
- Text
- URL

```
[14]: # Getting Top 10 positive news articles
pos_news = crisis_news_df.sort_values(by="compound", ascending=False)
pos_news = pos_news.head(10)
pos_news
```

```
[14]:
```

	compound	neg	neu	normalized	pos	source \
date						
2019-07-13	0.8831	0.073	0.689	1	0.238	The New York Times
2019-07-01	0.8481	0.027	0.769	1	0.204	Business Insider
2019-07-08	0.8360	0.000	0.809	1	0.191	Ozy.com
2019-07-08	0.8360	0.000	0.809	1	0.191	Ozy.com
2019-07-01	0.8240	0.000	0.807	1	0.193	Marketwatch.com
2019-07-08	0.7650	0.049	0.751	1	0.200	Marketwatch.com
2019-06-30	0.7184	0.000	0.878	1	0.122	Reuters
2019-07-18	0.7050	0.027	0.868	1	0.105	Daily Mail
2019-07-20	0.6956	0.149	0.595	1	0.256	The Hill
2019-07-24	0.6908	0.044	0.822	1	0.134	TechCrunch

text \

```
date
2019-07-13 Let me start with what might seem like a trivi...
2019-07-01 The drafters of the Declaration of Independenc...
2019-07-08 This is an OZY Special Briefing, an extension ...
2019-07-08 This is an OZY Special Briefing, an extension ...
2019-07-01 LONDON (Project Syndicate) The U.S. economy ha...
2019-07-08 The hunt for yield is making parts of the U.S...
2019-06-30 LONDON (Reuters) - Deutsche Bank plans to hire...
2019-07-18 After years of stability, Britain is facing pr...
2019-07-20 How many nations have already embraced sociali...
2019-07-24 Its the absolute best economy the United State...
```

url

```
date
2019-07-13 https://www.nytimes.com/2019/07/13/opinion/gol...
2019-07-01 https://www.businessinsider.com/50-maps-that-e...
2019-07-08 https://www.ozy.com/need-to-know/special-brief...
2019-07-08 https://www.ozy.com/need-to-know/deutsche-bank...
2019-07-01 https://www.marketwatch.com/story/old-age-does...
2019-07-08 https://www.marketwatch.com/story/yield-gap-be...
2019-06-30 https://www.reuters.com/article/us-deutsche-ba...
2019-07-18 https://www.dailymail.co.uk/news/article-72627...
2019-07-20 https://thehill.com/opinion/finance/453913-has...
2019-07-24 http://techcrunch.com/2019/07/24/why-do-media-...
```

```
[15]: # Create a table with hvplot
pos_news_table = pos_news.hvplot.table(
    columns=["date", "source", "text", "url"], width=500
)
pos_news_table
```

```
[15]: :Table    [date,source,text,url]
```

```
[16]: # Getting Top 10 negative news articles
neg_news = crisis_news_df.sort_values(by="compound", ascending=True)
neg_news = neg_news.head(10)
neg_news
```

```
[16]:
```

	date	compound	neg	neu	normalized	pos	source \
	2019-07-22	-0.9343	0.293	0.707	-1	0.000	Medium.com
	2019-07-10	-0.8783	0.198	0.802	-1	0.000	Marketwatch.com
	2019-07-12	-0.8779	0.206	0.794	-1	0.000	Marketwatch.com
	2019-07-13	-0.8658	0.251	0.749	-1	0.000	Hypebeast.com
	2019-06-25	-0.8625	0.206	0.794	-1	0.000	Politico
	2019-07-22	-0.8573	0.170	0.830	-1	0.000	Washingtonexaminer.com
	2019-07-22	-0.8442	0.214	0.786	-1	0.000	National Review
	2019-07-09	-0.8402	0.191	0.809	-1	0.000	Dailysignal.com
	2019-07-08	-0.7351	0.137	0.863	-1	0.000	TechCrunch
	2019-07-21	-0.6705	0.247	0.636	-1	0.117	Youtube.com

```
text \
```

	date	text
	2019-07-22	I warned about an economic crash years before ...
	2019-07-10	Bianca, a flight attendant living in Denver, h...
	2019-07-12	A stock-market index of small caps is at its w...
	2019-07-13	R. Kelly was arrested on several charges last ...
	2019-06-25	Michael Kruse is a senior staff writer for Pol...
	2019-07-22	Have you heard the news? Presidential candidat...
	2019-07-22	Senator Elizabeth Warren (D., Mass.) at a town...
	2019-07-09	Most\r\nof the public was outraged when, durin...
	2019-07-08	Technology has been used to manage regulatory ...
	2019-07-21	AbstractSpreadsheets are one of the most widel...

```
url
```

	date	url
	2019-07-22	https://medium.com/@teamwarren/the-coming-econ...
	2019-07-10	https://www.marketwatch.com/story/this-flight-...
	2019-07-12	https://www.marketwatch.com/story/this-stock-m...
	2019-07-13	https://hypebeast.com/2019/7/r-kelly-federal-r...
	2019-06-25	https://www.politico.com/magazine/story/2019/0...
	2019-07-22	https://www.washingtonexaminer.com/opinion/eli...
	2019-07-22	https://www.nationalreview.com/news/elizabeth-...


```

2019-07-09 https://www.dailysignal.com/2019/07/09/senate-...
2019-07-08 http://techcrunch.com/2019/07/08/the-startups-...
2019-07-21 https://www.youtube.com/watch?v=GyWKxFxyyRQ

```

```

[17]: # Create a table with hvplot
neg_news_table = neg_news.hvplot.table(
    columns=["date", "source", "text", "url"], width=500
)
neg_news_table

```

```

[17]: :Table [date,source,text,url]

```

1.1.7 Creating the Sentiment Distribution by News Article's Source

In this section, use hvPlot to create a bar chart that presents the distribution of negative, neutral and positive news according to the normalized score; the results should be grouped by source.

```

[18]: source_sentiment_chart_df = (
    crisis_news_df[["normalized", "source", "text"]]
    .groupby(["normalized", "source"])
    .count()
)
source_sentiment_chart_df.rename(
    index={-1: "Negative", 0: "Neutral", 1: "Positive"}, inplace=True
)
source_sentiment_chart_df

```

```

[18]:
normalized source text
Negative Aol.com 1
BBC News 1
Boingboing.net 1
Business Insider 6
Dailysignal.com 1
Fastcompany.com 2
Forbes.com 3
Gothamist.com 1
Huffpost.com 1
Hypebeast.com 1
Indianexpress.com 1
Marketwatch.com 6
Medium.com 1
Moneycontrol.com 1
Mywallst.com 1
National Review 1
Phys.org 1
Politico 2
Project-syndicate.org 3
Salon.com 1

```

	Seattletimes.com	1
	TechCrunch	1
	Thinkprogress.org	1
	Typepad.com	1
	Washingtonexaminer.com	2
	Youtube.com	1
Neutral	BBC News	3
	Digg.com	1
	Forbes.com	1
	Indianexpress.com	1
...	...	
	The Economist	1
	Theatlantic.com	1
Positive	Autoblog.comhttps	1
	Business Insider	6
	CNN	1
	Daily Mail	1
	Fool.com	2
	Forbes.com	3
	Hbr.org	1
	Huffpost.com	1
	Jalopnik.com	1
	Marketwatch.com	6
	Npr.org	1
	Ozy.com	2
	Politico	1
	Project-syndicate.org	2
	Psfk.com	1
	Qz.com	1
	Reuters	2
	Ritholtz.com	1
	Slate.com	1
	TechCrunch	1
	The Globe And Mail	2
	The Hill	1
	The New York Times	1
	Theatlantic.com	1
	USA Today	1
	Vice News	1
	Vox.com	1
	Yahoo.com	2

[62 rows x 1 columns]

```
[19]: source_sentiment_chart = source_sentiment_chart_df.hvplot.bar(
      xlabel="Sentiment",
      ylabel="Number of News",
```

```

        title="Sentiment Distribution by News Article's Source",
        height=450,
        width=1000,
        rot=90,
    )
    source_sentiment_chart

```

```
[19]: :Bars [normalized,source] (text)
```

1.1.8 Creating the Word Clouds

In this section you will create two word clouds, one using the bag-of-words method and other using TF-IDF.

Bag-of-Words' Word Cloud Use the CountVectorizer module from sklearn to create a word cloud with the top 20 words with the highest counting. Save the DataFrame with the top 20 words as a CSV file named top_words_data.csv for future use on the dashboard creation.

```
[20]: # Creating the CountVectorizer instance defining the stop words in English to
      ↪be ignored
      vectorizer = CountVectorizer(stop_words="english")
```

```
[21]: # Getting the tokenization and occurrence counting
      X_words = vectorizer.fit_transform(crisis_news_df["text"])
```

```
[22]: # Retrieve unique words list
      words = vectorizer.get_feature_names()

      # Get the last 100 word (just as a sample)
      print(words[-100:])
```

```

['underdogs', 'underestimate', 'unemployment', 'unfolded', 'unfortunately',
'uninterrupted', 'union', 'united', 'university', 'unleashed', 'unless',
'unlocking', 'unnoticed', 'unsettling', 'unsplash', 'unthinkable', 'unveiled',
'upright', 'urban', 'usa', 'use', 'used', 'users', 'valley', 'various', 'vast',
've', 'venezuela', 'verge', 'versus', 'vessels', 'veteran', 'vice', 'victims',
'video', 'view', 'viewers', 'views', 'visit', 'vladimir', 'volatile', 'vriend',
'vucci', 'wage', 'wager', 'wall', 'wants', 'warned', 'warning', 'warren',
'warrenhas', 'washington', 'wasn', 'wasnt', 'way', 'weak', 'weakest', 'wealth',
'week', 'weekly', 'wel', 'wh', 'whats', 'white', 'widely', 'widespread',
'williamson', 'win', 'window', 'wing', 'winning', 'winters', 'wisdom', 'wise',
'witnesses', 'wobble', 'wonder', 'wondering', 'wont', 'words', 'work',
'working', 'world', 'wouldnt', 'writer', 'writers', 'written', 'yang', 'year',
'years', 'yield', 'yielded', 'yielding', 'yields', 'york', 'youtuber', 'zombie',
'zombies', 'zone', 'zurich']

```

```
[23]: # Getting the bag of words as DataFrame
      words_df = pd.DataFrame(
```

```

    list(zip(words, np.ravel(X_words.sum(axis=0)))), columns=["Word",
→"Word_Count"]
)

# Sorting words by 'Word_Count' in descending order
words_df.sort_values(by="Word_Count", ascending=False, inplace=True)

```

```

[24]: # Get top 20 words with the highest counting
top_words = words_df.head(20)
top_words

```

```

[24]:
      Word  Word_Count
343   chars           99
850     li            57
601  financial         26
424   crisis          21
1494  warren           19
35    2008            18
1283  share           18
1442    ul            17
527  elizabeth        17
1431   trump          16
1100 president        13
232   bank            12
976    new            12
515  economic         12
1054   plan           11
1533   year           11
650   global          11
905   market          10
520  economy          10
499  donald           10

```

```

[25]: # Save the top words DataFrame
file_path = Path("Data/top_words_data.csv")
top_words.to_csv(file_path)

```

```

[26]: # Create a string list of terms to generate the bag-of-words word cloud
terms_list = str(top_words["Word"].tolist())

```

```

[27]: # Create the bag-of-words word cloud
wordcloud = WordCloud(colormap="RdYlBu").generate(terms_list)
fig_bow_cloud = plt.figure()
plot_bow_cloud = plt.imshow(wordcloud)
plot_bow_cloud = plt.axis("off")
fontdict = {"fontsize": 20, "fontweight": "bold"}
plot_bow_cloud = plt.title("Bag-of-Words Wordcloud", fontdict=fontdict)
plot_bow_cloud = plt.show()
plt.close(fig_bow_cloud)

```

Bag-of-Words Wordcloud



TF-IDF Wordcloud Use the TfidfVectorizer module from sklearn to create a word cloud with the top 20 words with the highest frequency. Save the DataFrame with the top 20 words as a CSV file named top_wors_tfidf_data.csv for future use on the dashboard creation.

```
[28]: # Getting the TF-IDF
tfidf_vectorizer = TfidfVectorizer(stop_words="english")
X_tfidf = tfidf_vectorizer.fit_transform(crisis_news_df["text"])
```

```
[29]: # Retrieve words list from corpus
words_tfidf = tfidf_vectorizer.get_feature_names()

# Get the last 100 word (just as a sample)
print(words_tfidf[-100:])
```

```
['underdogs', 'underestimate', 'unemployment', 'unfolded', 'unfortunately',
'uninterrupted', 'union', 'united', 'university', 'unleashed', 'unless',
'unlocking', 'unnoticed', 'unsettling', 'unsplash', 'unthinkable', 'unveiled',
'upright', 'urban', 'usa', 'use', 'used', 'users', 'valley', 'various', 'vast',
've', 'venezuela', 'verge', 'versus', 'vessels', 'veteran', 'vice', 'victims',
'video', 'view', 'viewers', 'views', 'visit', 'vladimir', 'volatile', 'vriend',
'vucci', 'wage', 'wager', 'wall', 'wants', 'warned', 'warning', 'warren',
'warrenhas', 'washington', 'wasn', 'wasnt', 'way', 'weak', 'weakest', 'wealth',
'week', 'weekly', 'wel', 'wh', 'whats', 'white', 'widely', 'widespread',
'williamson', 'win', 'window', 'wing', 'winning', 'winters', 'wisdom', 'wise',
'witnesses', 'wobble', 'wonder', 'wondering', 'wont', 'words', 'work',
'working', 'world', 'wouldnt', 'writer', 'writers', 'written', 'yang', 'year',
'years', 'yield', 'yielded', 'yielding', 'yields', 'york', 'youtuber', 'zombie',
'zombies', 'zone', 'zurich']
```

```
[30]: # Creating a DataFrame Representation of the TF-IDF results
words_tfidf_df = pd.DataFrame(
    list(zip(words_tfidf, np.ravel(X_tfidf.sum(axis=0)))), columns=["Word",
    ↪ "Frequency"]
)

# Sorting words by 'Frequency' in descending order
words_tfidf_df = words_tfidf_df.sort_values(by=["Frequency"], ascending=False)
```

```
[31]: # Get 20 top words
top_words_tfidf = words_tfidf_df.head(20)
top_words_tfidf
```

```
[31]:
```


	Word	Frequency
850	li	6.039592
343	chars	4.139945
601	financial	2.740431
1494	warren	2.601068
527	elizabeth	2.274471
424	crisis	2.269101
1283	share	2.160880
35	2008	2.060891
1431	trump	1.966178
1442	ul	1.859178
232	bank	1.697471
1054	plan	1.674686
1100	president	1.638754
1533	year	1.630204
515	economic	1.617877
650	global	1.617172
976	new	1.563557
520	economy	1.494965
905	market	1.402006
499	donald	1.352434

```
[32]: # Save the top words TF-IDF DataFrame
file_path = Path("Data/top_wors_tfidf_data.csv")
top_words_tfidf.to_csv(file_path)
```

```
[33]: # Create a string list of terms to generate the tf-idf word cloud
terms_list_tfidf = str(top_words_tfidf["Word"].tolist())
```

```
[34]: # Create the tf-idf word cloud
wordcloud_tfidf = WordCloud(colormap="RdYlBu").generate(terms_list_tfidf)
fig_tfidf_cloud = plt.figure()
plot_tfidf_cloud = plt.imshow(wordcloud_tfidf)
plot_tfidf_cloud = plt.axis("off")
fontdict = {"fontsize": 20, "fontweight": "bold"}
plot_tfidf_cloud = plt.title("TF-IDF Wordcloud", fontdict=fontdict)
```

```
plot_tfidf_cloud = plt.show()
plt.close(fig_tfidf_cloud)
```



1.2 Challenge: Radar Chart with Tone Analysis

In this challenge section, you have to use Plotly Express and IBM Watson Tone Analyzer to create a radar chart presenting the tone of all the news articles that you retrieved.

Refer to the [polar coordinates chart demo](#) and the [Plotly Express reference documentation](#) to learn more about how to create this chart.

[35]: *# Get the Tone Analyzer API Key and URL*

```
tone_api = os.getenv("tone_api")
tone_url = os.getenv("tone_url")
```

[36]: *# Initialize Tone Analyzer Client*

```
tone_analyzer = ToneAnalyzerV3(version="2017-09-21", iam_apikey=tone_api,
    url=tone_url)
```

In order to create the radar chart, you need to score the tone of each article and retrieve the document_tone. Create a function named `get_tone(text,url)` that will receive two parameters and will get the tone score for a particular article.

- `text` the content of the article.
- `url` the URL pointing to the article.

The `get_tone()` function will use the `tone()` method from the `ToneAnalyzerV3` module to score the article's tone. Remember that for each document (or text), the `tone()` method of IBM Watson Tone Analyzer [scores one or more overall document tones](#), you can also get an empty result if no tone were scored; this function should return a dictionary with the first document tone's score with the following structure:

- score refers to the first tone from the document_tone.
- tone_id refers to the tone_id from the first tone.
- tone_name refers to the tone_name from the first tone.
- text the text passed as parameter.
- url the URL passed as parameter.

This is an example of the function's return value:

```
{'score': 0.616581,
'tone_id': 'sadness',
'tone_name': 'Sadness',
'text': '\nMore than a decade since the global economic meltdown of 2008
devastated lives across the world, no one who caused the crisis has
been held responsible.\n\n"The 2008 financial crisis displayed what
the world now identifies as financial contagion," says Philip J Baker,
the former managing partner of a US-based \nhedge fund that collapsed
during the financial crisis.\n\nDespite this, "zero Wall Street chief
executives have been to prison, even though there is today absolutely
no doubt that Wall Street executives and politicians \nwere complicit
in creating the crisis," he says. \n\nBaker was among the few
relatively smaller players imprisoned for the part they played.\n\n
In July 2009, he was arrested in Germany and extradited to the
United States where he faced federal court on charges of fraud and
financial crimes.\n\nHe pled guilty and was sentenced to 20 years
in prison for costing some 900 investors about $294mn worldwide.
He served eight years in jail and is now on \nparole and advocates
against financial crime.\n',
'url': 'https://www.aljazeera.com/programmes/specialseries/2019/06/men-stole-world-2008-financ
```

```
[37]: # Create a function to analyze the text's tone with the 'tone()' method of IBM
      ↪ Watson Tone Analyzer.
def get_tone(text, url):
    try:
        tone_analysis = tone_analyzer.tone(
            {"text": text},
            content_type="application/json",
            sentences=False,
            content_language="en",
            accept_language="en",
        ).get_result()

        result = tone_analysis["document_tone"]["tones"][0]
        result["text"] = text
        result["url"] = url
        return result
    except:
        pass
```


Create a DataFrame with the tone scoring from all the news articles. Use an empty list to create a the DataFrame's structure and a for-loop to iterate across all the news to score their tone using the get_tone() function.

```
[38]: # Create an empty list to create the DataFrame's structure
articles_tone_data = []

# Iterate across all the news articles to score their tone.
print(f"Analyzing tone from {crisis_news_df.shape[0]} articles...")
for index, row in crisis_news_df.iterrows():
    try:
        print("*", end="")
        # Get news article's tone
        article_tone = get_tone(row["text"], row["url"])
        if article_tone: # Validate if a tone where found
            articles_tone_data.append(article_tone)
    except:
        pass
print("\nDone :-)")
```

Analyzing tone from 100 articles...

Done :-)

```
[39]: # Create the DataFrame containing the news articles and their tone scoring
      → results.
articles_tone_df = pd.DataFrame.from_dict(articles_tone_data)
articles_tone_df.score = articles_tone_df.score.round(2)
articles_tone_df.head()
```

```
[39]:
```

	score	text	tone_id \
0	0.66	Share to facebookShare to twi...	tentative
1	0.56	Michael Kruse is a senior staff writer for Pol...	sadness
2	0.50	Andrew Yang didnt set out to become a test pre...	tentative
3	0.53	A WELL-KNOWN stockmarket sell signal is a comp...	tentative
4	0.69	A new high for the U.S. stock market before a ...	joy

	tone_name	url
0	Tentative	https://www.forbes.com/sites/mayrarodriguezval...
1	Sadness	https://www.politico.com/magazine/story/2019/0...
2	Tentative	https://slate.com/news-and-politics/2019/06/an...
3	Tentative	https://www.economist.com/finance-and-economic...
4	Joy	https://www.marketwatch.com/story/a-china-trad...

Save the DataFrame as a CSV file named tone_data.csv for further use on the dashboard creation.

```
[40]: articles_tone_df.to_csv(Path("Data/tone_data.csv"), index=False)
```

Create a radar chart using the scatter_polar() method from Plotly Express as follows:

- Use the score column for the r and color parameters.
- Use the tone_name column for the theta parameter.
- Use the url column for the hover_data parameter.
- Define a title for the chart.

```
[41]: tone_radar = px.bar_polar(
    articles_tone_df,
    r="score",
    theta="tone_name",
    color="score",
    hover_data=["url"],
    title="News Articles Tone Analysis",
)
tone_radar.show()
```

News Articles Tone Analysis

